



DTE-2502: NEURAL NETWORKS

MODULE02: GRADIENT DESCENT

BASIC DEFINITIONS

Kalyan Ram Ayyalasomayajula, email: kay001@uit.no

1. GRADIENT DESCENT

Basic Concept for single variable function:

- Function : $f(x)$ that we want to minimize
- Gradient : $\nabla f(x) = \text{derivative/slope at point } x$
- Update Rule : $x_{new} = x_{old} - \alpha \cdot \nabla f(x)$
- Learning Rate : α (step size parameter)

Extending to multiple variable function

- Function : $f(x_1, x_2, \dots, x_n)$
- Gradient : $\nabla f = [\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n}]$

Algorithm 1 Gradient Descent Algorithm

Require: Training data, cost function $Q(w)$

Ensure: Optimized parameters (weights) w^*

- 1: Initialize parameters w (randomly or zeros)
 - 2: Set learning rate α
 - 3: **repeat**
 - 4: Calculate gradient: $\nabla Q(w)$
 - 5: Update parameters: $w \leftarrow w - \alpha \cdot \nabla Q(w)$
 - 6: Check convergence criteria
 - 7: **until** convergence
 - 8: **return** optimized parameters w
-

Traditional gradient descent (batch gradient descent) uses average gradient computed over the entire dataset for each update. This is a very simplified version (and not the entire picture) of what happens in gradient descent, but good enough to understand.

- Pros : Stable convergence, guaranteed to converge for convex functions
- Cons : Slow for large datasets, memory intensive

2. STOCHASTIC GRADIENT DESCENT

Minimization of average loss over the training data

$$Q(w) = \frac{1}{l} \sum_{i=1}^l \mathcal{L}_i(w) \rightarrow \min_w$$

Input: dataset X^l , learning rate μ , parameter λ

Output: weights $w = (w_{jh}, w_{hm})$

Initialization

Set all the weights w to small random numbers
Evaluate the objective function $Q(w)$

do


select x_i from X^l

compute the loss function $\mathcal{L}_i(w)$

gradient step $w := w - \mu \mathcal{L}'_i(w)$

update the objective function $Q := \lambda \mathcal{L}_i + (1 - \lambda) Q$

until Q and/or w converges



3. MINI-BATCH GRADIENT DESCENT

Mini-batch gradient descent uses average of gradients computed over a mini-batch (small set) of dataset for each update.

- Pros : Balance between batch-GD and SGD
- Cons : Need to tune optimal batch size

4. GRADIENT DESCENT OPTIMIZERS

In all the below approaches $\nabla f(\theta)$ is the gradient w.r.t the parameter vector θ .

DEFINITION 1: Momentum.

Key idea: As the descent direction is fixed add historical vector to accelerate convergence $v_t = v_{t-1} + \epsilon \nabla f(\theta_{t-1})$

DEFINITION 2: Nesterov momentum (NAG).

Key idea: In addition to the descent direction "Look ahead" by computing gradient at anticipated future position. $v_t = v_{t-1} + \epsilon \nabla f(\theta_{t-1} - \mu v_{t-1})$

DEFINITION 3: AdaGrad.

Key idea: Adaptive learning rates - larger updates for infrequent parameters, smaller updates for frequent ones.

$$v_t = v_{t-1} + g^2, \text{ where } g = \nabla f(\theta_{t-1})$$

$$\theta_{t+1} = \theta_t - \frac{\mu}{\sqrt{v_t} + \epsilon} g$$

- Pros

- No manual learning rate tuning per parameter

- Great for sparse data
- Larger updates for rare features
- Cons
 - Learning rate decreases too aggressively, may stop learning

DEFINITION 4: RMSprop.

Key Idea : Fix AdaGrad's vanishing learning rate using exponential moving average.

$$v_t = \beta v_{t-1} + (1 - \beta)g^2, \beta \approx 1$$

$$\theta_{t+1} = \theta_t - \frac{\mu}{\sqrt{v_t} + \epsilon} g$$

- Pros
 - Maintains adaptive learning rates without vanishing
 - Works well for non-stationary objectives
- Cons
 - Sensitive to choice of μ

DEFINITION 5: Adam (Adaptive Moment Estimation).

Key Idea : Combine momentum + RMSprop with bias correction.

$$v_t = \beta_1 v_{t-1} + (1 - \beta_1)g; \hat{v}_t = \frac{v_t}{1 - \beta_1^t}$$

$$s_t = \beta_2 s_{t-1} + (1 - \beta_2)g^2; \hat{s}_t = \frac{s_t}{1 - \beta_2^t}$$

$$\theta_{t+1} = \theta_t - \frac{\mu \hat{v}_t}{\sqrt{\hat{s}_t} + \epsilon} g$$

5. PRACTICAL ISSUES IN NEURAL NETWORKS TRAINING

- Overfitting
 - Regularization
 - Dropout
- Vanishing or exploding gradients
 - Adaptive learning rate
 - Batch normalization
- Local optima
 - Pre-training